

A STORYBOARD IN 10 FRAMES

What counts as good, and how to check it

*Graded on vibes.
Shipped on hope.*



Evaluation compares behaviour against an expectation written down before the run. Every faster, cheaper check approximates that written judgment.

“Is it good?”

Good at what?



*moves without saying
which part moved*

Does it parse?

*fields present,
types right*

Is it supported?

*by the source
it cites*

Does it decline?


*what it's
meant to*

Right register?

*the one the product
committed to*

Each property fails differently and is checked by different means.
One blended score hides which of them moved.

Decide what counts before you look.



Goalposts, relocated post-kick.

Evaluation compares observed behaviour against an expectation that was written down before the run.

Four checks, one judgment.

1

Deterministic assertions

exact, cheap, narrow

⋮

2

Heuristics

a tripwire that catches the obvious

⋮

3

Model-based graders

judgment at volume, same probabilistic kind

⋮


4

Human review

the one the other three stand in for

Every cheaper check approximates a judgment someone made first. The judgment doesn't disappear when it's automated; it gets encoded.


The model marked itself.



Model-graded. Model-approved. A+.

Model-based graders are one of four ways to check, each an approximation of the written judgment. The expensive one, people, is the one the other three stand in for.

Keep the bad outputs.



change something

rerun the same set


see what moved

only your change moved

~~fresh samples~~ → **the same set**

It isn't an estimate of production traffic. It's failures that happened at least once, kept so their return gets noticed instead of reported by a user.

A metric should change a decision.



z z z

Tracked everything. Decided nothing.

A retained metric should say what is measured, on what population, over what period, how it was produced, and which decision changes if it moves.

A number needs a decision.

1

What is measured

the property, named

⋮

2

On what population

whose runs, which cases

⋮

3

Over what period

from when, to when

⋮

4

How it was produced

the method behind the number

⋮


5

Which decision changes

if none would, drop the metric

A figure quoted without its population and method hardens into a fact, and the conditions it was true under stop travelling with it.



Everyone agreed it felt better.



Five thumbs up. Zero written criteria.

Evaluation only becomes possible once someone writes down what good means for this task. Every faster check approximates that written judgment.

Write down what “worse” looks like.



before the change, not after

Before the next change, describe behaviour you wouldn't accept, concretely enough that someone else could look at an output and agree it happened. Everything else checks that.

Read the whole explanation

[pruningmypothos.com/systems/
evaluation-is-a-human-problem/](https://pruningmypothos.com/systems/evaluation-is-a-human-problem/)

