

A STORYBOARD IN 11 FRAMES

# From agent intent to governed execution

*any day now.**It can ask. It can't sign.*

THE GATE

THE TOOL

THE RECORD

A model can ask for an action. Something else decides whether that ask becomes a real effect, and then whether it actually worked.



# Four things happen before money moves.



## 1. Check the request

*a real operation? right types? things that exist?*



## 2. Decide if it's permitted

*for this caller, in this context, right now*



## 3. Run it

*with whatever limits apply*




## 4. Check what happened

*did the change actually land?*

A sequence, not an architecture. Small systems collapse it into a few lines of code; it still helps to know which line is doing which.



# Each step fails in its own way.



**Valid, but not allowed**

*the gate said no. Rightly, or not.*



**Allowed, but it didn't run**

*the tool errored, or timed out*




**It ran. Nothing changed.**

*the most confident failure of all*

A request can be valid but not allowed, allowed but fail to run, and run successfully without producing the effect anyone wanted.

# Four things people mistake for permission.



~~Logged in~~  
tells me who.  
Not what.



~~Well formed~~  
So is "delete  
production."



~~"97% sure"~~  
That's a mood,  
not a permit.



**A human said yes**  
A way to answer.  
Not the question.

A rule in a prompt is a request. A rule at the gate is a check.



# Perfect paperwork. Still no.



*Every box ticked. Every box irrelevant.*

Authorization asks whether an operation is permitted, which is a different question from whether the request is well-formed.

# “Not sure” means several things.



*yes/no only: block the useful, or wave the rest through*

## ALLOW


permitted, and the evidence holds

*leaves: a receipt*

## ASK

a person should decide

*leaves: a question, and why*



## DENY

not permitted, or too risky

*leaves: a reason code*

## DEFER

thin evidence, high stakes

*leaves: a handoff note*

*what moves the needle most:  
how hard is it to undo?*

*the bar for ALLOW rises →*




undo: one click



undo: an apology email




# It said “delivered.”



*Delivered. To number 12. You live at 21.*

A successful tool call is not proof that the intended effect occurred. Checking what actually happened is its own step.

# “Ask a human” has to reach one.



Who reviews it, what they see, and whether they are positioned to say no is its own design problem: human judgment, the next stage of the map.



# The tool says done. The record disagrees.




*...is it, though?*

~~reports done~~ → **shows done**

A 200 tells you what the interface said. If what you wanted was a change of state, go and read the state.

# “It broke.” Which part?



“...and then the  
agent just...”

incident review, as storytelling



**Policy denied it**  
*was the rule right?*



**The tool errored**  
*fix it, or retry it*



**Verification failed**  
*the tool fibbed, or you  
checked the wrong thing*



**Evidence was missing**  
*go and get it*




**Over time or budget**  
*raise the limit, or don't*

Each tray points somewhere different. You can't sort a failure at a boundary you never made explicit.



# The model still varies. Your effects needn't.



The checks don't make the model deterministic. They give you a dependable answer about one candidate action: permitted or not, every time.

## Read the whole explanation

[pruningmypothos.com/systems/  
from-agent-intent-to-governed-execution/](https://pruningmypothos.com/systems/from-agent-intent-to-governed-execution/)

