


A STORYBOARD IN 6 FRAMES



# When is an AI system ready for real use?

*Deployed. Delighted.  
Disowned.*




Deployment means it's running. Readiness means the evidence and controls justify this exposure: these users, doing this job, with these consequences when it's wrong.

# Not ready in general. Ready for something.



*a team reviewing it*



*customers acting on it*



*same system. different exposure.*


Name the exposure precisely: which users, doing what, with what riding on the result. That turns an unanswerable question into one with an answer.

# The demo proved it's possible.



## The demo

*one user, one path,  
little at stake*





## Real use

*the cases nobody  
demoed*

Readiness asks about conditions the demonstration may never have exercised. Evidence that's silent about a case isn't a pass.

# Set the bar before you look.



*fitting the rule to the results*

Define the release criterion against an evaluation set that fits the use and its failure costs, before exposure. Chosen afterwards, it can be fitted to the results.

# Rollback isn't universal.



*some effects  
don't come back.*



## Stop it

*before more  
happens*

## Reverse it

*not a sent message  
or a payment*

## Correct it

*after the fact*


## Contain it

*limit how far  
it travels*

The question isn't whether a rollback procedure exists. It's what can be stopped, reversed, corrected or contained.

# Three answers, in writing.

- ☐ Evidence against the failures that matter
- ☐ What limits the damage when it's incomplete
- ☐ Who decides, and who answers



*named, not assumed*

If those can't be answered, the finding isn't that the system is unready. It's that nobody can say yet, which is a different problem and a more fixable one.

## Read the whole explanation

[pruningmypothos.com/systems/  
from-prompt-to-production/](https://pruningmypothos.com/systems/from-prompt-to-production/)

