



A STORYBOARD IN 9 FRAMES

What an AI model actually is

*A file of numbers.
Worshipped accordingly.*




built around it, by people:




An AI model is a file of numbers that turns input into a ranking over what comes next. Retrieval, tools, memory, checks and policy are separate software built around it.

It's a file. Of numbers.



0.0213 -1.447
0.8811 0.004
-0.339 1.217
0.5006 -0.071
1.9031 0.318
-0.772 0.064
...
(billions more)

fixed after training




there's no such row

*no single number
is a fact*

Training set the numbers. After training they are fixed, and every run uses the same ones. They encode patterns useful for predicting text, spread across billions of values.

Same question. Different answer.



Second opinion. Same doctor.

The model picks from a ranking rather than always taking the top choice, so the same input can come back worded differently on another run. Expect variation.

One run, one piece at a time.



Cut the text into tokens

each piece becomes a number



Score every token

the parameters rank what could come next



Pick one

not always the top one, so wording varies




Add it, run again

that loop is what "generation" means

This describes language models specifically. Image models, speech models and classifiers are models too, and they don't work this way.

Likely isn't true.



How likely it scored

the model's own ranking

~~How confident it sounds~~


a property of the wording

Whether it's true

nothing in the loop checks this

A confident tone is a pattern in the training text, not a report on the model's reliability. If you want a truth check, something else has to perform it.

It doesn't remember you.



Every conversation is a first date.

Memory is stored state kept by the software around the model. The model itself carries nothing from one run to the next.

It didn't look that up.



separate software, each fixable



~~It looked it up~~
retrieval did.


~~It refused~~
policy did.

~~It remembered~~
stored state did.

It ranked
what comes next.
That part was it.

When the output is wrong, the model is one suspect among several, and often not the right one.


It sounds sure. It learned to.



Confidence is a writing style.

A confident tone is a pattern in the text the model was trained on, not a report on how reliable this answer is. If you want a truth check, something else has to do it.

Hand it the facts. Don't ask it to hold them.



Give it the facts rather than asking it to hold them. Expect variation between runs, and when something goes wrong, ask which part failed before blaming the model.

Read the whole explanation

[pruningmypothos.com/systems/
what-an-ai-model-actually-is/](https://pruningmypothos.com/systems/what-an-ai-model-actually-is/)

