



A STORYBOARD IN 7 FRAMES

What an AI model actually is

*A file of numbers.
Worshipped accordingly.*




built around it, by people:




An AI model is a file of numbers that turns input into a ranking over what comes next. Retrieval, tools, memory, checks and policy are separate software built around it.

It's a file. Of numbers.



0.0213 -1.447
0.8811 0.004
-0.339 1.217
0.5006 -0.071
1.9031 0.318
-0.772 0.064
...
(billions more)

fixed after training



there's no such row

*no single number
is a fact*

Training set the numbers. After training they are fixed, and every run uses the same ones. They encode patterns useful for predicting text, spread across billions of values.

One run, one piece at a time.



Cut the text into tokens

each piece becomes a number



Score every token

the parameters rank what could come next



Pick one

not always the top one, so wording varies




Add it, run again

that loop is what "generation" means

This describes language models specifically. Image models, speech models and classifiers are models too, and they don't work this way.

Very likely. Also wrong.



How likely it scored

the model's own ranking

~~How confident it sounds~~


a property of the wording

Whether it's true

nothing in the loop checks this

A confident tone is a pattern in the training text, not a report on the model's reliability. If you want a truth check, something else has to perform it.

It sounds sure. It learned to.



Confidence is a writing style.

A confident tone is a pattern in the text the model was trained on. It says nothing about how reliable this answer is; a truth check has to come from somewhere else.

It didn't look that up.



separate software, each fixable



~~It looked it up~~

retrieval did.

~~It refused~~

policy did.

~~It remembered~~


stored state did.

It ranked

*what comes next.
That part was it.*

When the output is wrong, the model is one suspect among several, and often not the right one.

Hand it the facts. Don't ask it to hold them.



Give it the facts rather than asking it to hold them. Expect variation between runs, and when something goes wrong, ask which part failed before blaming the model.

Read the whole explanation

[pruningmypothos.com/systems/
what-an-ai-model-actually-is/](https://pruningmypothos.com/systems/what-an-ai-model-actually-is/)

